

Data Science Indy: News

July 2019

Data Science News

Google announced the availability of preconfigured container images with deep learning frameworks and tools.

<https://www.forbes.com/sites/janakirammsv/2019/07/01/google-just-made-machine-learning-more-accessible-and-portable-with-containers/#6ad4b0894dcc>

Facebook AI Research announced the open-source release of a deep-learning recommendation model, DLRM, that achieves state-of-the-art accuracy in generating personalized recommendations.

<https://www.infoq.com/news/2019/07/facebook-recommendation-model/>

Tensorflow has released a pruning API:

<https://medium.com/tensorflow/tensorflow-model-optimization-toolkit-pruning-api-42cac9157a6a>

This is in addition to other work Google has been doing on improving performance, cost to serve, and carbon footprint

<https://ai.googleblog.com/2019/05/efficientnet-improving-accuracy-and.html?m=1>

Data Science News

"Computer science curriculums should spend more time emphasizing how deep learning works and especially the limitations of today's correlative pattern encoders."

<https://www.forbes.com/sites/kalevleetaru/2019/07/10/why-is-there-so-little-understanding-of-deep-learnings-limitations/#a3c3f97b96fd>

Contrary to McKinsey and others stating that AI may kill or displace millions of workers, ZipRecruiter asserted AI created about three times as many jobs as it took away in 2018

<https://www.fastcompany.com/90369739/is-ai-killing-jobs-actually-it-added-3x-more-in-2018>

Data Science News

Author pulls program which attempts to create inappropriate photos of unwitting image subjects

<https://arstechnica.com/tech-policy/2019/06/author-pulls-software-that-used-deep-learning-to-virtually-undress-women/>

Deep learning has a terrible carbon foot print: a single AI model can emit as much carbon as five cars in their lifetimes

<http://arxiv.org/abs/1906.02243>

Facebook stole Bill Gate's voice

<https://www.theverge.com/2019/6/10/18659897/ai-voice-clone-bill-gates-facebook-melnet-speech-generation>

Links or Discussion

Notebook with several top Kaggle implementations:

<https://www.kaggle.com/sudalairajkumar/winning-solutions-of-kaggle-competitions>

PyTorch has its own Model Zoo: <https://pytorch.org/hub>

AWS offers one competing approach to data quality:

<https://aws.amazon.com/blogs/big-data/test-data-quality-at-scale-with-deequ/>

Word embeddings may not be as biased as once thought:

<https://twitter.com/rikvannoord/status/1132933236756238341>